

(Due: April 27)

Project Description

The purpose of this project is to write a distributed application using the Remote Procedure Call (RPC) protocol and the CUDA toolkit. In this project you are expected to use the ONC RPC Protocol to distribute computation from a Linux workstation to a Linux workstation with the lowest load.

This project can be divided into two parts. The client program running on a workstation provides the interface to the users. The servers running on the pre-selected four Linux workstations provide the machine load reporting service and another two computational services. You are required to follow the eight steps of RPC distributed application implementation discussed in the class and use the protocol compiler `rpcgen` to implement the client/server programs.

The client program keeps a list of four available Linux workstations in our lab. It firstly invokes a remote procedure `getload()` to collect the load average (over the last 1 minute) from each workstation in the list and selects the workstation with the lowest load average. Then, based on the command line option (`-gpu` or `-lst`), the client invokes the computation function `findmax_gpu()` or `update_lst()` on the selected workstation, gets the result back, and prints out the result.

Two examples are shown below. The first example, with the command line arguments `-gpu N M S`, finds the largest element of an array which has 2^N (e.g. 2^{20}) elements. The elements of the array have been initialized using the exponential distribution with the mean value M (e.g. 5) and with the initial seed S (e.g. 17). The computation is executed on the smallest load machine chopin with its GPU. In the second example, the client calls the remote function `update_lst()` on the machine bach which converts each value F in the linked list by using the formula $\sqrt[4]{F} * 20.0$ and returns the updated linked list back to the client.

```
haydn % ldshr -gpu 20 5 17
arthur: 2.30  bach: 3.3  chopin: 0.50  degas: 1.27
(executed on chopin)
75.161147
```

```
haydn % ldshr -lst 5 16 81 25 49 625
arthur: 1.30  bach: 0.50 chopin: 2.27  degas: 9.99
(executed on bach)
29.9 40.0 60.0 44.7 52.9 100.0
```

The server provides three services (i.e functions). To get the load average (over the last 1 minute) on Linux, you can call the system function `getloadavg()`. To find the largest element of an array, the function `findmax_gpu()` firstly initializes the very large array using the values N , M , and S passed from the client. It then launches the kernel function on GPU to find the largest element in the array and returns the value back to the client. You are required to modify the code in `reduction_kernel.cu`

discussed in the class to implement this. Similarly, the function `update_lst()` gets a linked list of doubles from the client and then updates the values in the list.

Note that you should run the server program on each of the four Linux workstations before you start the client.

Turnin

Each group (at most two students) needs to submit this project using the following turnin command:

```
turnin -c cis620s -p proj4 report.pdf makefile ldshr.x ldshr.c \  
ldshr_svc_proc.c findmax.cu
```

Your report should include the description of your design/implementation, experiences in debugging/testing, project status(works or not), etc. The cover page should contain your picture(s), name(s) and the login id you used to turnin the project. Start on time and good luck. If you have any questions, send e-mail to j.sang@csuohio.edu.